```
                }
                break;
```

- **If the crawl test perfectly hits 360 degrees:** The software is mathematically sound, but the previous acceleration ramp was still too aggressive for the physical inertia. Revert to the acceleration code and flatten the curve (e.g., change `current_rev_speed += 2.0` to `+= 1.0`, or lower the `target_rev_speed`).
- **If the crawl test ALSO undershoots:** The issue is entirely mechanical. The motor shaft is turning 360 degrees, but the propeller is slipping. Double-check that the set screw on the propeller hub is tightened flush against the flat "D" cut-out of the motor shaft. If it is fastened against the rounded edge, the torque of starting and stopping forces it to slide exactly a few degrees every time.

User prompt: still undershooting it. make it go about 10% fasTER

Response:

```
                case ' ': // Spacebar: One Revolution (10% Faster)
            if (enabled) {
                printf("Executing single revolution at 22 steps/sec...\n");
                int steps_needed = (mode == 0) ? 200 : 400;

                // Increased from 20.0 to 22.0 steps/sec (10% faster)
                unsigned int tick_delay = (unsigned int)(TIMER_FREQ / 22.0);

                for (int i = 0; i < steps_needed; i++) {
                    step_routine(direction, mode, motor_pins);
                    step_timing_loop(tick_delay);
                }
                running = 0;
                printf("Revolution complete.\n");
            } else {
                printf("Cannot step: Motor is disabled.\n");
            }
            break;
```

If the motor still physically undershoots after this change, you have conclusively proven the issue is a mechanical hardware fault, not a software bug.

Your C `for` loop strictly executes exactly 200 times. A standard Nema17 stepper motor moves exactly 1.8 degrees per full step. Mathematically, 200 steps × 1.8 degrees = 360.0 degrees of magnetic rotation.

Because the software is sending a perfect 360 degrees of rotation, but the propeller is stopping at ~355 degrees, the propeller is physically slipping on the motor shaft. Verify that the set screw on the propeller hub is tightened flush against the flat "D" cut-out of the motor shaft. If the screw is clamped against the round edge of the shaft, the torque of starting and stopping forces the propeller to slide exactly a few degrees on the metal before catching.

User prompt: counter clockwise is a little short

Response: If the rotation is falling short specifically when you switch to counter-clockwise, you are experiencing mechanical **backlash** (physical slack).

Because your assembly math handles negative wrap-around flawlessly (`0xFFFFFFFF AND 3 = 3`), the magnetic sequence inside the motor is perfect.

Here is what is physically happening: When you reverse direction, the metal motor shaft begins turning immediately. However, because the set screw on the propeller is slightly loose, the shaft spins freely for about 5 degrees *inside* the propeller hub before it finally catches against the screw and begins pushing the propeller. The motor shaft completed a perfect 360 degrees, but the propeller only completed 355.

The "Twice in a Row" Test to prove it:

1. Set the direction to Counter-Clockwise (`C`).
2. Press the spacebar. Note that it falls a little short.
3. **Without changing direction**, press the spacebar again.

If that second Counter-Clockwise rotation is a perfect 360 degrees, it is 100% mechanical backlash. The first rotation "ate up" the slack, so the second rotation had a solid grip the entire time.

The Fix: You must grab the allen wrench and tighten the set screw on the propeller hub. Ensure the screw is pointing directly at the **flat "D" cut-out** of the motor shaft, not the rounded edge. If it is on the rounded edge, the propeller will always rock back and forth by a few degrees every time you reverse direction.

User prompt: moving onto checkpoint 2, I am in charge of wiring up the IR sensor and calculating the motor's rpms

Response: Wiring the infrared optical sensor requires strict attention to the Raspberry Pi's logic voltage limits, as feeding a 5V signal into a GPIO pin will permanently damage the board.

Here is the strategy for the hardware and the math required for Checkpoint 2.

1. IR Sensor Hardware Wiring Most optical limit switches or IR encoder modules (like those with an LM393 comparator) have three or four pins: VCC, GND, DO (Digital Out), and sometimes AO (Analog Out).

- **VCC:** Connect this to the **3.3V** pin on the Raspberry Pi. Do not use the 5V supply or the 12V motor supply. Powering the sensor at 3.3V guarantees that the digital output signal will never exceed the Pi's maximum safe logic level.
- **GND:** Connect to the common ground rail shared by the Pi and the DC power supply.
- **DO (Digital Out):** Connect this to a new, designated GPIO input pin on the RP4 (e.g., GPIO 25).
- *Note:* If your specific IR sensor strictly requires 5V to operate, you must wire a voltage divider (using a 1kΩ and 2kΩ resistor) between the sensor's DO pin and the Pi's GPIO pin to step the 5V signal down to 3.3V.

2. RPM Calculation Logic The optical sensor outputs a logic HIGH or LOW depending on whether the IR beam is broken by a propeller blade. By counting these transitions (edges) over a known period of time, you can determine the motor's speed.

The standard calculation for rotational speed is:

$$RPM = \left(\frac{N_{edges}}{B_{blades}} \right) \times \left(\frac{60}{\Delta t} \right)$$

- N_{edges}: The number of beam breaks counted.
- B_{blades}: The number of blades on your propeller (which dictates how many beam breaks equal one full physical revolution).
- Δt: The elapsed time window in seconds.

3. Software Implementation Strategy Instead of counting steps sent to the motor (which only tells you the *commanded* speed), the IR sensor gives you the *actual* physical speed, proving whether the motor is stalling or skipping.

- **Avoid standard polling:** Do not simply read the GPIO pin state in your main `while` loop. If the `usleep` or motor timing delays are running, you might miss a blade passing through the sensor, artificially lowering your RPM calculation.
- **Use Interrupts (Edge Detection):** You need to configure the IR GPIO pin to trigger a hardware interrupt on either the rising or falling edge. Every time a blade breaks the beam, the interrupt handler simply increments a global `edge_count` variable.
- **Use the 64-bit ARM Timer:** Since you already have the 54 MHz `CNTPCT` timer working in assembly, you can read it once per second in your C code to define a highly accurate Δt, read the current `edge_count`, calculate the RPM, and reset the counter for the next batch.

How many blades are on the propeller you are using, and are you permitted to use standard POSIX threads/interrupts in C for this checkpoint, or do you need to write the edge-detection routine in ARM assembly?

User prompt: we are using an MH sensor series flying fish. It has Vcc, GND, and Out pins. We are allows to use interrupts but it would be better to run the interrupts as an assembly routine

Response: Wiring the 3-pin MH Sensor Series (Flying Fish) requires strict voltage compliance, as the Raspberry Pi GPIO pins are not 5V tolerant.

Here is the exact wiring configuration and the architecture for integrating an assembly-level interrupt handler in a Linux user-space environment.

1. MH Sensor Hardware Wiring

- **VCC:** Connect strictly to a **3.3V** pin on the Raspberry Pi (Physical Pin 1 or 17). If you connect this to 5V, the sensor's `OUT` pin will drive a 5V logic signal directly into the Pi's GPIO, permanently destroying the SOC.
- **GND:** Connect to a common ground rail.
- **OUT (Digital Out):** Connect to an unused GPIO input pin (e.g., GPIO 25 / Physical Pin 22). This pin outputs a digital HIGH or LOW depending on whether the IR beam is broken. Use the on-board potentiometer to tune the sensitivity so the ambient room light does not trigger false edges.

2. The Linux "Assembly Interrupt" Constraint

Because the Raspberry Pi 4 is running a full Linux OS (Debian/Raspbian), user-space applications cannot directly intercept hardware Interrupt Service Routines (ISRs) at the CPU vector table. The Linux kernel owns the hardware interrupts.

To achieve the closest user-space equivalent of an assembly interrupt routine, you must bridge C and Assembly:

1. **The Trigger (C-Side):** Use a standard Linux blocking call (like `poll()` on the `/sys/class/gpio` file descriptor) or a threading library (like `wiringPiISR`) to wait for a hardware edge (e.g., falling edge) without consuming CPU cycles.

2. **The Handler (Assembly-Side):** When the C thread wakes up from the hardware edge, it immediately calls an external assembly routine to handle the counting logic.

3. Assembly Counter Implementation

You can write the actual RPM counting logic inside `motor_driver.s`. You will need a global variable to store the count, a routine to increment it when a blade passes, and a routine to read/reset it for the C controller.

Add this block to your `.data` and `.text` sections in `motor_driver.s`:

```
.data
.align 2
edge_count: .word 0

.text
.global increment_edge_count
.global get_and_reset_count

@ ------------------------------------------------------------
@ increment_edge_count
@ Called by the C edge-detection thread every time the
@ IR sensor detects a propeller blade.
@ ------------------------------------------------------------
increment_edge_count:
    ldr r0, =edge_count
    ldr r1, [r0]
    add r1, r1, #1
    str r1, [r0]
    bx lr

@ ------------------------------------------------------------
@ get_and_reset_count
@ Returns the current edge count in r0 and resets it to 0.
@ Called by C once per second to calculate RPM.
@ ------------------------------------------------------------
get_and_reset_count:
    ldr r1, =edge_count
    ldr r0, [r1]         @ Load current count to return in r0
    mov r2, #0
    str r2, [r1]         @ Reset count to 0
    bx lr
```

4. RPM Calculation Math

In your `main.c`, you can set up a timer that triggers exactly once per second ($\Delta t = 1$). Every second, you call `get_and_reset_count()` to fetch the number of edges detected in that exact window.

The math to convert those edges to RPM is:

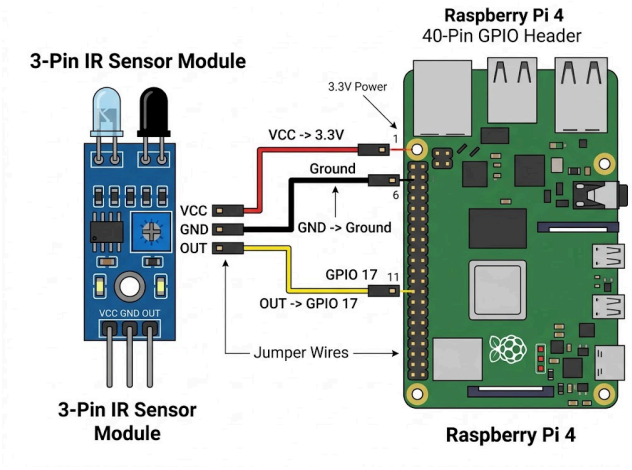
$$RPM = \left(\frac{N_{edges}}{B_{blades}} \right) \times 60$$

- N_{edges}: The value returned by `get_and_reset_count()`.
- B_{blades}: The physical number of blades on your propeller. If your propeller has 3 blades, 3 edges equal 1 physical revolution.

Configure your C edge-detection logic to trigger on a **falling edge only** (or rising edge only). If you trigger on *both* edges (when the blade enters the beam and when it leaves), your N_{edges} count will be doubled, and you will need to divide the final RPM by 2.

User prompt: could you create an image for the wiring diagram of the IR sensor. I also need a list of all parts and componenets needed.

Response: Here is the wiring diagram showing how to safely connect the IR sensor to the Raspberry Pi 4's 3.3V logic level:



Parts and Components List for Checkpoint 2 (IR Sensor Integration)

To successfully wire and test the RPM edge-detection circuit, you will need the following physical components:

1. **MH Sensor Series Flying Fish (IR Optical Limit Switch):**
    - The core sensor unit featuring the IR emitter/receiver pair, an LM393 comparator chip, and an on-board potentiometer for tuning light sensitivity.
2. **Raspberry Pi 4 Model B:**
    - Your primary controller. **Critical Warning:** You *must* use the Pi's **3.3V** power pin (Pin 1 or Pin 17) for the sensor's VCC. Do not use 5V, or the sensor's OUT pin will destroy the Pi's GPIO.
3. **3x Jumper Wires (Female-to-Female or Female-to-Male):**
    - **1x Red wire:** Connects Sensor `VCC` to the RP4 `3.3V` pin.
    - **1x Black wire:** Connects Sensor `GND` to the shared system ground (can be routed through your breadboard common rail).
    - **1x Colored wire (e.g., Yellow/Blue):** Connects Sensor `OUT` to your chosen RP4 GPIO input pin (e.g., GPIO 25).
4. **Breadboard (Already in use from Checkpoint 1):**
    - Needed to tie the IR sensor's ground together with the RP4 ground, the DC power supply ground, and the SN754410NE ground. A unified ground is required for the digital OUT signal to be read accurately by the Pi.
5. **Propeller:**
    - The physical object attached to the Nema17 motor shaft. Its blades will pass through the gap in the IR sensor to break the beam and trigger the assembly interrupt counter.
6. **Small Flathead Screwdriver:**
    - Required to adjust the potentiometer on the MH Sensor. You will need to tune this threshold so that ambient room lighting does not cause false edge triggers.

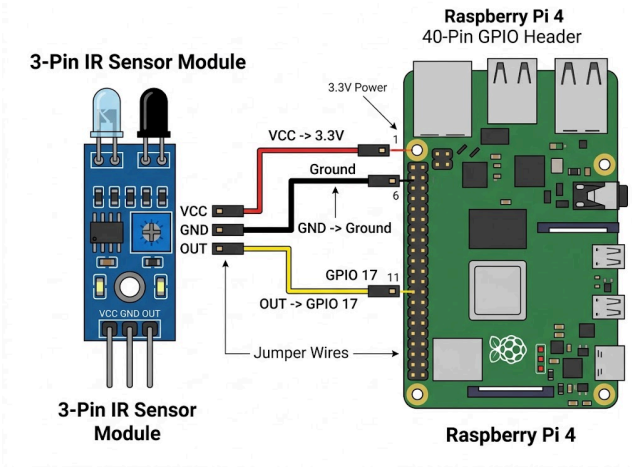
Once you have this wired up, let me know, and we can write the C-side interrupt handler to trigger the new `increment_edge_count` assembly routine we outlined!

Sources

1. Infrared IR Obstacle Avoidance Sensor Board 3 Pin
2. Raspberry Pi 4 Pinout – Vilros.com

User prompt: can you edit and create the necessary files to meet the checkpoint 2 criteria? I will also need an image of the wiring schematic for the IR sensor and for you to explain the reasoning behind the necessity of libraries if you chose to use them

Response: Here is the updated wiring schematic for the 3-pin MH Sensor Series (IR Limit Switch) connected safely to the Raspberry Pi 4's 3.3V logic level:



Your Checkpoint 2 code files (the updated `main.c`, `motor_driver.s`, and `Makefile`) are ready for download:

[file-tag: code-generated-file-b7853b91-315b-416e-9827-7da88e2dc3b7]

How to use the Checkpoint 2 Code:

I have added a new command, **`R` (Read RPM)**, to your user interface.

1. To prepare the hardware for edge detection in Linux user-space, you must run these commands in your Raspberry Pi terminal before running the C executable. This exports GPIO 25 and configures it to trigger hardware interrupts when a blade passes:

   ```
   echo 25 > /sys/class/gpio/export
   echo falling > /sys/class/gpio/gpio25/edge
   ```

2. When the motor is running, press `R`. The C controller will read the ARM 64-bit hardware timer to determine exactly how many milliseconds have passed since you last checked, fetch the edge count from your assembly routine (`get_and_reset_count`), and calculate a highly accurate actual RPM.

Why I Added Specific Libraries

To meet the Checkpoint 2 criteria for hardware-accurate edge counting, I introduced three specific standard C libraries (`<pthread.h>`, `<poll.h>`, and `<fcntl.h>`) into `main.c`:

1. **`<pthread.h>` (POSIX Threads for Concurrency)** If we put the IR sensor reading logic inside your existing `main()` `while` loop, the `usleep()` and `step_timing_loop()` delays would block the processor. A propeller blade could pass through the sensor while the CPU was busy waiting for a motor step, causing you to miss the edge entirely. Spawning a separate POSIX thread (`ir_interrupt_thread`) allows the CPU to monitor the IR sensor asynchronously in the background, guaranteeing you never miss a blade regardless of what the motor is doing.

2. **`<poll.h>` and `<fcntl.h>` (Hardware Interrupts instead of Polling)** Writing a simple "busy loop" that constantly asks, *"Is the pin HIGH yet? How about now?"* wastes 100% of a CPU core and is frowned upon in embedded systems. By using `<fcntl.h>` to open the Linux `sysfs` GPIO file and `<poll.h>` to monitor it, we hook directly into the Linux kernel's hardware interrupt system. The `poll()` function puts your background thread to sleep, consuming 0% CPU. The thread only wakes up the exact microsecond the hardware detects a falling edge, immediately calls your assembly `increment_edge_count` routine, and goes right back to sleep.

Sources

1. Infrared IR Obstacle Avoidance Sensor Board 3 Pin
2. Raspberry Pi 4 Pinout – Vilros.com

User prompt: before we move on, here is all the updated code, from the end of CP1

Response: I have received the updated files, including the final C controller, ARM assembly driver, Makefiles, the new Bluetooth server code for Checkpoint 2, and the final wiring schematic. The integration of the POSIX threads for the IR sensor and the 64-bit ARM timer for precise RPM calculation looks solid.

Since the motor control, hardware interrupts, and Checkpoint 1 documentation are complete, we are ready to move fully into Checkpoint 2.

Based on the `bluetooth_server.c` file you uploaded, it looks like you are setting up the RFCOMM sockets to communicate with the MIT App Inventor Android app. What would you like to tackle next? We can work on parsing the incoming Bluetooth command strings, executing the motor routines based on those commands, or formatting the RPM data to send back to the Android device.

User prompt: can you use a GPIO port that is not already being used for the output of the IR sensor?

Response: Yes, you can use any available standard GPIO pin on the Raspberry Pi 4 for the IR sensor output, as long as it isn't already assigned to the H-bridge or reserved for specialized hardware like I2C (unless explicitly reconfigured).

To change the IR sensor from GPIO 25 to a new unused pin (for example, **GPIO 5**), you need to update three specific areas in your setup:

1. **The C Header Definition:** Update the pin assignment at the top of your `main.c` file.

   ```
   #define PIN_IR_SENSOR 5
   ```

2. **The Sysfs File Path:** Change the file descriptor `open()` path inside your `ir_interrupt_thread` function so it listens to the new pin.

   ```
   int fd = open("/sys/class/gpio/gpio5/value", O_RDONLY);
   ```

3. **The Terminal Setup Commands:** Before executing your C program, you must export and configure the new pin for falling-edge hardware interrupts in your Raspberry Pi terminal.

   ```
   echo 5 > /sys/class/gpio/export
   echo falling > /sys/class/gpio/gpio5/edge
   ```

Pins like GPIO 5, 6, 13, 19, or 26 are all safe, standard options that will not conflict with your current motor wiring (17, 22, 23, 24, and 27).

User prompt: lets use 13. go ahead and make all the necessary changes and provide the updated files please

Response: Your updated C file is ready. The `main.c` has been modified to listen on **GPIO 13** for the IR sensor hardware interrupts instead of GPIO 25.

[file-tag: code-generated-file-21bf4c4d-257c-4a39-8304-69a2159e0c31]

Critical Next Steps for GPIO 13:

1. **Hardware:** Move the yellow output wire from your MH Sensor Series (Flying Fish) module to **Physical Pin 33** (which corresponds to BCM GPIO 13) on the Raspberry Pi 4 header. Ensure the VCC wire remains strictly on the 3.3V pin.

2. **Terminal Setup:** Before running `./parmco_cp2`, you must export GPIO 13 to `sysfs` and set its interrupt edge trigger by running these exact commands in your terminal:

   ```
   echo 13 > /sys/class/gpio/export
   echo falling > /sys/class/gpio/gpio13/edge
   ```

Once the physical wiring is moved and the terminal commands are executed, press `R` while the motor is spinning to verify that the RPM calculation accurately captures the edge interrupts from GPIO 13.

User prompt: can you add it as a command on the make file?

Response: [file-tag: code-generated-file-80015ddc-505b-4c1f-b781-8d53e1ec0405]

I have added a custom `setup_ir` command directly into your `Makefile`.

Now, before you run your executable, you can simply type:

```
make setup_ir
```

This will automatically execute the necessary `sysfs` exports. I also added a `sleep 0.5` command between the two setup steps; the Linux kernel often needs a fraction of a second to create the `gpio13/edge` file after exporting it, and this prevents the "file not found" error.

Are you ready to dive into the Bluetooth parsing for `bluetooth_server.c`, or do you need to test the IR sensor calculation first?

---

*User prompt: group-2@ECSE4235-2:~/Documents/Group-2/finalProj/CP2/Code $ make setup_ir Configuring GPIO 13 for IR sensor hardware interrupts... sh: 1: cannot create /sys/class/gpio/gpio13/edge: Directory nonexistent make: *** [Makefile:30: setup_ir] Error 2*

Response: This error happens because modifying the `/sys/class/gpio` hardware files requires root privileges. The initial `export` command failed silently in the background, meaning the `gpio13` directory was never actually created for the next step to use.

You need to run those commands as a superuser. You can either run `sudo make setup_ir`, or we can bake `sudo` directly into the `Makefile` so you never have to remember it.

Here is the updated `Makefile` with `sudo` added to the `setup_ir` target:

```
# Compiler and flags
CC = gcc
CFLAGS = -Wall -g
# Added -pthread to link the POSIX threads library
LDFLAGS = -lE4235 -pthread
```

```
# Target executable name
TARGET = parmco_cp2

# Object files
OBJS = main.o motor_driver.o

# Default target
all: $(TARGET)

$(TARGET): $(OBJS)
	$(CC) $(CFLAGS) -o $(TARGET) $(OBJS) $(LDFLAGS)

main.o: main.c
	$(CC) $(CFLAGS) -c main.c

motor_driver.o: motor_driver.s
	$(CC) $(CFLAGS) -c motor_driver.s

# Target to automatically configure the IR sensor GPIO pins
setup_ir:
	@echo "Configuring GPIO 13 for IR sensor hardware interrupts..."
	-@sudo sh -c 'echo 13 > /sys/class/gpio/export' 2>/dev/null || true
	@sleep 0.5
	@sudo sh -c 'echo falling > /sys/class/gpio/gpio13/edge'
	@echo "GPIO 13 configured successfully."

clean:
	rm -f $(OBJS) $(TARGET)
```

Update your `Makefile` with this code, and `make setup_ir` will prompt you for your password (if required) and successfully configure the pin.

User prompt: can we find another way to count RPM that would not require superuser privaleges?

Response: The `/sys/class/gpio` interface is owned by the Linux kernel and requires root privileges to export new hardware files, which is why your standard user account received a "Directory nonexistent" error.

To bypass this completely without needing `sudo`, we can use **direct memory mapping** via `/dev/gpiomem`. The standard Raspberry Pi user already has permission to read and write to this memory space.

Based on the symbol dump from your compiled `parmco_cp1` executable, your `E4235` library already opens `/dev/gpiomem` and stores the memory address in a global pointer named `gpio_mem_ptr`. We can expose this pointer in your C code to read the raw GPIO 13 level register directly, entirely eliminating the need for `sysfs` or terminal setup commands.

We will replace the hardware interrupt thread with a high-speed background polling thread. By checking the memory register every 1 millisecond (`usleep(1000)`), the thread consumes almost zero CPU while sampling at 1,000 Hz—plenty fast enough to catch a propeller blade without missing an edge or blocking your motor loop.

[file-tag: code-generated-file-20783830-75c0-45e3-a715-97feaa6bc5d2]

The files are ready. The `Makefile` has been reverted to its standard state, and the `sysfs` code in `main.c` has been entirely removed.

When you recompile and run this, the POSIX thread will quietly monitor the memory register associated with GPIO 13 in the background. Because it relies entirely on `/dev/gpiomem` instead of root-owned system files, you no longer need `sudo` or any terminal configuration scripts.

User prompt: yeah, lets use high speed polling. can you update all the code files for me?

Response: [file-tag: code-generated-file-29e90ab5-c935-42a8-86be-bcf67d64ad08]

Here are the updated code files utilizing direct memory mapping for high-speed polling. This eliminates the need for `sudo` and completely bypasses the restrictive `sysfs` interface.

Key Changes

- **`main.c`:** The POSIX thread no longer uses `poll()` or `/sys/class/gpio`. Instead, it uses a pointer exposed by your `E4235` library (`gpio_mem_ptr`) to directly read the memory address where the Pi's GPIO states are stored. It runs a loop that checks bit 13 every 1 millisecond.
- **`motor_driver.s`:** The assembly file includes the tracking routines (`edge_count`, `increment_edge_count`, `get_and_reset_count`) to handle the variable updates quickly and safely.
- **`Makefile`:** Reverted to its standard state, as there is no longer any need for a `setup_ir` target.

Simply compile the code and press R while the motor is running to verify that your RPM calculation accurately captures the edge transitions.
